

nnsreg: Reportable segment slopes for nonlinear relationships in Stata

Eric A. Booth

Texas 2036

Austin, TX

eric.a.booth@gmail.com

Abstract. The `nnsreg` command brings univariate Nonlinear Nonparametric Statistics (NNS) regression (Viole and Nawrocki 2013) to Stata and returns the fitted line segments as ordinary estimation results. The command first uses recursive partial-moment partitions to place regression points and linear-spline knots. Its default `method(connect)` joins the NNS regression points, while `method(ols)` holds the NNS-derived knots fixed and estimates the spline coefficients by least squares. Both versions store the knot locations and range-specific slopes, so users can turn a nonlinear shape into reproducible statements about where the relationship is flat, steep, or changing and can compare that partition-based description with a kernel or series fit. In simulations, the two `npregress` estimators recover smooth curves more accurately on average. In the clustered simulation, the NNS least-squares spline nearly matches the most accurate smoother and has much lower between-replication variability. On the motorcycle-impact data, the series fit has the lowest mean cross-validation error, with the NNS least-squares spline and kernel smoother close behind. These results support using `nnsreg` to report range-specific slopes alongside a smoother. Comparable predictive error appears only in the clustered simulation and motorcycle benchmark.

Keywords: st0000, nnsreg, nonparametric regression, partial moment, linear spline, knot, postestimation

1 Introduction

Where does a relationship change? The applied analyst who asks this question, in education finance, program evaluation, or health-services research, usually needs both a picture of the relationship and a few reproducible statements about where it is flat, steep, or changing. Test outcomes that stop improving after spending passes a level, infrastructure that deteriorates slowly until a threshold age and quickly afterward, and risk that accelerates only in the tail each raise the same descriptive question: where does the fitted curve bend, and how steep is it within each range? A single regression slope cannot answer them, because it averages over the whole range of the data. A flexible smoother such as `npregress kernel` (StataCorp 2025) can display the shape, but its output is a curve, and the analyst must do additional work, with `margins` or by hand, to turn it into a small set of reportable slope statements. Making a flexible fit easy to

. This is a preprint. The article has been submitted to the *Stata Journal* and is currently under review.

report is itself a recognized Stata problem: `xb1c` tabulates and plots flexible covariate fits (Orsini and Greenland 2011), and `marginscontplot` plots the marginal effects of continuous predictors (Royston 2013). `nnsreg` works in that tradition and returns a different object: named ranges and slopes as ordinary estimation results.

Stata already offers several routes to a nonlinear fit, and `nnsreg` sits among them rather than replacing them. `mk spline` builds a linear spline but asks the analyst to set the knots, either by hand or at sample percentiles, which respond to the marginal distribution of the regressor but not to its joint pattern with the outcome; placing knots adaptively is a long-standing idea in the spline literature (Friedman 1991). Polynomials and fractional polynomials (Royston and Altman 1994) bend smoothly but do not name the places where the slope changes. `npregress series` fits a spline or polynomial basis with a data-driven number of terms and posts estimation results, and `npregress kernel` traces a smooth mean. The community-contributed `binsreg` partitions the regressor’s support into data-driven quantile bins and fits within each bin, with formal inference (Cattaneo et al. 2024, 2025).¹ And when the question is the location of one or a few breaks, with inference on where they lie, formal segmented-regression and joinpoint estimators (Muggeo 2003; Kim et al. 2000), and Stata’s own `threshold` command, answer it directly; `nnsreg` is an exploratory, descriptive complement to those, not a substitute. What `nnsreg` adds is a particular way of placing the segment boundaries, the NNS partition of the data, paired with coefficients named for the first slope and each later slope change.

`nnsreg` takes its way of placing boundaries from an established method. Nonlinear Nonparametric Statistics (NNS), developed by Viole and Nawrocki (2013), builds regression fits from partial moments, which are averages of deviations below and above a target value (Viole and Nawrocki 2012a,b; Nawrocki and Viole 2024). The NNS regression algorithm recursively partitions the scatterplot into quadrants, computes a representative regression point inside each partition, and connects the points with line segments. Vinod and Viole (2018b) compare this approach with kernel regression from the R `np` package (Hayfield and Racine 2008) and report three practical advantages: the method requires no bandwidth choice, the segment slopes estimate derivatives directly, and the partitions adapt to data whose support is clustered rather than spread evenly. They also report that in-sample fit is similar for the two approaches, with differences they describe as minor. A companion paper (Vinod and Viole 2018a) describes the NNS regression points as analogues of linear-spline knots. The reference implementation is the NNS package for R (Viole 2026).

This article contributes a Stata implementation and reporting workflow for an existing estimator. `nnsreg` brings the univariate core of NNS regression to Stata as a deliberate subset of the R package (section 2.3 lists the differences). It also makes explicit the two parts of the regression fit. The NNS partition chooses the regression

1. I do not benchmark `binsreg` on accuracy in what follows because it targets a different object: a binned conditional-mean function with uniform confidence bands, rather than a spline with named, reportable segment slopes. Its bin selection is tuned for inference on the mean function, not for minimizing curve-recovery error, so an accuracy row would put tools built for different jobs side by side.

points and knot locations; the user then either connects those points or estimates a least-squares regression on the resulting linear-spline basis. In both cases, the first-segment slope and each subsequent change in slope are named coefficients in `e(b)`, the knot locations are stored in `e(knots)`, and `predict`, `estimates store`, `lincom`, and `coefplot` (Jann 2014) work as they do after `regress`. Other Stata commands make a nonstandard estimate usable the same way: Newson’s `somersd` (Newson 2006), for one, posts rank statistics as a fitted model so that `lincom` can difference them. With `nnsreg`, the analyst tests a claim like “the fitted slope is flat below age 35 and about 1.2 points per year at the oldest ages” with two `lincom` commands on the stored coefficients, instead of reading it off a graph by eye.

`nnsreg` is descriptive: a bend in a fitted curve is a feature of the conditional mean in the sample, not a causal threshold, and causal interpretation still requires a research design. The command is also not a general replacement for `npregress kernel`; in my simulations the kernel smoother recovers smooth curves more accurately on average. The evidence here supports considering `nnsreg` when the empirical support is clustered or gapped, or when the deliverable is a segment-level summary with standard postestimation tools.

Section 2 explains how the method works. Section 3 documents the syntax, options, and stored results. Section 4 compares the command with linear regression and `npregress kernel` and `npregress series` on three simulated designs with known truth, including a Monte Carlo study across 100 replications. Section 5 repeats the comparison on the real motorcycle-impact data, scored by cross-validation. Section 6 offers guidance on when to use which tool, and section 7 summarizes the findings and the command’s limits.

2 The NNS regression approach

2.1 Partial moments and recursive partitioning

The data consist of pairs (x_i, y_i) , and the target is the conditional mean of y as a function of x without committing to a functional form. The NNS approach starts from partial moments, which summarize a variable’s behavior on one side of a target value t . One-sided moments of this kind are an established tool with a five-decade history in risk measurement (Bawa 1975; Fishburn 1977); the NNS contribution is to pair targets on both variables of a bivariate sample and to build a regression method on the resulting quadrants. The lower and upper partial moments of degree r are

$$\text{LPM}_r(t; X) = \frac{1}{N} \sum_{i=1}^N \max(0, t - X_i)^r$$

$$\text{UPM}_r(t; X) = \frac{1}{N} \sum_{i=1}^N \max(0, X_i - t)^r$$

At degree zero, the positive part raised to zero is defined as an indicator of a strictly positive deviation. The two expressions therefore become $N^{-1} \sum_i I(X_i < t)$ and $N^{-1} \sum_i I(X_i > t)$; observations tied at t contribute to neither moment. The recursive partition assigns ties to the lower quadrant so every observation belongs to one group. Partial moments thus generalize the empirical distribution function (Viole and Nawrocki 2012a). The useful fact for regression is that a pair of targets, one for X and one for Y , divides the scatterplot into four quadrants, and those quadrants describe the local joint behavior of the two variables (Viole and Nawrocki 2012b).

The algorithm applies this idea recursively. Starting from the full sample, **nnsreg** computes a center for X and a center for Y (the mean by default, the median with **noise(median)**), splits the observations into the four quadrants around that center, and then repeats the split inside every quadrant that still holds more than **minobs()** observations (eight by default). Splitting therefore reaches more finely into dense regions than into sparse ones. The **order()** option controls the depth: order one uses the global regression point, and each higher order permits another generation of partitions. A higher order therefore permits a more finely segmented fit; the default of three balances resolution against noise across the designs studied here (section 4.6). With **partition(xonly)**, the command instead splits only along X , which produces vertical bands; this variant is useful as a sensitivity check and for applications where partitioning on the outcome is undesirable.

2.2 From regression points to a linear spline

After partitioning stops, **nnsreg** collapses each final group to a single regression point, its center in (x, y) space. It then adds two boundary points, one at the smallest and one at the largest observed x . Each boundary y is a count-weighted average of two local linear fits evaluated at that boundary: one fit over the observations out to the nearest regression point, and one out to the midpoint between the boundary and that point; with fewer than six observations in the first window, or no spread in x in the second, the command falls back to the mean of the distinct y values observed at the boundary. It adds one central point, placed in x midway between the two middle regression points by rank, with y the mean of the observations that fall between them (or the average of those two points' y values when none do). It then sorts all points by x and merges any that share an x value by averaging their y ; there is no minimum-spacing rule, so two distinct but nearby x values remain separate knots. Finally it connects the ordered points. The connected segments are then expressed as a linear spline. With interior knots $k_1, \dots, k_m$ placed at the interior regression points, the fitted function is

$$\hat{f}(x) = \beta_0 + \beta_1 x + \sum_{j=1}^m \gamma_j (x - k_j)_+$$

where $(x - k_j)_+ = \max(0, x - k_j)$. The coefficient β_1 is the slope of the first segment, and each γ_j is the change in slope at knot k_j . This representation follows Vinod and Viole (2018a), who describe NNS regression points as playing the role that knots play in a linear spline; the $(x - k_j)_+$ basis itself is the standard truncated-line representation of

a linear spline in semiparametric regression (Ruppert et al. 2003). Because a piecewise-linear curve is a regression on a spline basis, **nnsreg** can post, store, and manipulate the fit with standard Stata tools.

nnsreg can interpolate or estimate coefficients on that basis. The default connected fit passes through the NNS regression points and is closest to the construction in the NNS literature. With **method(ols)**, the command keeps the NNS-derived knots but estimates the spline coefficients by ordinary least squares. This refinement extends the original construction by retaining only the NNS-derived knot locations. The connected fit summarizes the partition structure itself; the OLS fit uses the same segmentation but minimizes squared error given the knots, which in my simulations recovers the true curve more accurately (section 4).

The standard errors require careful interpretation. The OLS fit reports the usual least-squares standard errors, conditional on the data-derived knot locations. For the connected fit, **nnsreg** rescales the OLS spline covariance matrix by the connected fit's residual variance. Those intervals are heuristic references, not estimator-specific sampling intervals. Neither version accounts for selecting the knots from the same data. Inference that does account for an estimated break or kink location is available for segmented models with one or a few knots (Muggeo 2003; Hansen 2017), and extending it to the NNS partition is beyond this article. I therefore treat the OLS intervals as conditional descriptions and the connected-fit intervals as heuristic references, not as tests. Graphs, support counts, comparison fits, and the sensitivity tools in section 4.6 carry more weight.

2.3 Relationship to the R NNS package

The NNS package for R (Viole 2026) is the reference implementation of this method, and **nnsreg** deliberately implements a subset of it. Users moving between the two should know the differences. The R function **NNS.reg** supports multivariate regressors, classification-style prediction, and automatic selection of the partition depth based on an internal dependence measure interpreted as a signal-to-noise ratio (Vinod and Viole 2018a). **nnsreg** is univariate, uses a fixed user-chosen **order()**, and offers mean or median partition centers. In exchange, it integrates with Stata's estimation machinery: **e-class** results, **predict** in and out of the estimation sample, and compatibility with **estimates**, **lincom**, and **coefplot**. Users who need multivariate NNS regression or automated order selection should use the R package; users whose deliverable is a Stata workflow with reportable segment coefficients are the audience for **nnsreg**.²

2. Because the two implementations differ in these details, they agree in shape rather than term for term, and they scale the **order()** argument differently, so matching them requires matching the number of knots rather than the order. A fidelity check on the motorcycle data of section 5, distributed with the development repository, confirms that the fitted values track each other closely once matched this way.

3 The nnsreg command

3.1 Syntax

```
nnsreg depvar varname [if] [in] [, order(#) noise(mean | median)
    method(connect | ols) partition(xy | xonly) minobs(#) level(#)
    vce(vcetype) slopes noplot knotlines graph_opts(twoway_options)
    generate(newvar) ]
```

`vce(vcetype)` may be `robust` or `cluster clustvar` and is allowed only with `method(ols)`.

After estimation, `predict` follows the usual syntax:

```
predict [type] newvar [if] [in] [, xb residuals]
```

where `xb` (the default) computes fitted values from the stored spline and `residuals` computes the difference between the observed outcome and the fitted values; the two options are mutually exclusive. Because the spline basis is rebuilt from `e(knots)` at prediction time, `predict` also scores observations outside the estimation sample; the terminal segments extrapolate linearly, and `predict` prints a note whenever it scores points beyond the estimation range stored in `e(xmin)` and `e(xmax)`. `nnsreg` runs under Stata 16 or later.

3.2 Options

`order(#)` sets the partition depth, that is, how many rounds of recursive splitting are allowed. A higher order permits more line segments, which can follow curvature more closely but can also track noise. The default is `order(3)`. Section 4.6 shows how to read a small grid of orders as a sensitivity check.

`noise(mean | median)` sets the center used for partitioning and for the regression points. `noise(mean)`, the default, uses arithmetic means. `noise(median)` uses medians, which move less when a partition contains outlying or skewed values.

`method(connect | ols)` sets how spline coefficients are assigned. `method(connect)`, the default, interpolates the NNS regression points and is the construction closest to the NNS literature. `method(ols)` keeps the NNS-derived knots and chooses the spline coefficients that minimize squared error conditional on those knots.

`partition(xy | xonly)` sets the partition style. `partition(xy)`, the default, splits on both variables in quadrants, following the NNS construction. `partition(xonly)` splits only along the regressor, producing vertical bands for comparison with the default quadrant partition.

`minobs(#)` sets the smallest group the recursive partition keeps splitting: a group is split only if it holds more than `#` observations. The default is `minobs(8)`. Larger

values yield fewer, longer segments and make a useful sensitivity check alongside `order()`. The reported x ranges lie between adjacent regression-point knots, so a range can contain fewer than `#` observations, or none at all, even though every partition group obeys `minobs()`.

`level(#)` sets the confidence level for the reported intervals; the default is `level(95)` or as set by `set level`.

`vce(vctype)` requests a robust or cluster-robust covariance estimate for the least-squares stage; *vctype* may be `robust` or `cluster clustvar`. It is allowed only with `method(ols)`. The resulting standard errors remain conditional on the data-derived knots.

`slopes` reports a table of per-segment slopes, each the cumulative sum of the spline coefficients up to that segment, with delta-method standard errors. These are the quantities the `lincom` statements below compute by hand; the same table is stored in `e(segments)`.

`noplot` suppresses the default scatter-and-fit graph.

`knotlines` overlays vertical dashed lines at the interior knot locations on the default graph.

`graph_opts(twoway_options)` passes options through to the default graph.

`generate(newvar)` creates a variable holding the fitted values over the estimation sample.

3.3 Stored results

`nnsreg` stores the following in `e()`:

Scalars

<code>e(N)</code>	number of observations
<code>e(N_clust)</code>	number of clusters when <code>vce(cluster ...)</code> is specified
<code>e(r2)</code>	coefficient of determination, $1 - \sum e_i^2 / \sum (y_i - \bar{y})^2$
<code>e(r2.a)</code>	adjusted coefficient of determination
<code>e(rmse)</code>	root mean squared error, $(\sum e_i^2 / df_r)^{1/2}$
<code>e(rss)</code>	residual sum of squares
<code>e(mss)</code>	model sum of squares
<code>e(F)</code>	model F statistic; missing for <code>method(connect)</code>
<code>e(l1)</code>	log likelihood; missing for <code>method(connect)</code>
<code>e(l1.0)</code>	constant-only log likelihood; missing for <code>method(connect)</code>
<code>e(order)</code>	partition depth
<code>e(minobs)</code>	minimum group size for splitting
<code>e(df_m)</code>	model degrees of freedom
<code>e(df_r)</code>	residual degrees of freedom
<code>e(xmin)</code>	minimum of the regressor in the estimation sample
<code>e(xmax)</code>	maximum of the regressor in the estimation sample

Macros

<code>e(cmd)</code>	<code>nnsreg</code>
<code>e(cmdline)</code>	command as typed
<code>e(title)</code>	estimation-command title
<code>e(estat_cmd)</code>	<code>nnsreg_estat</code> , which rejects unsupported <code>estat</code> calls
<code>e(predict)</code>	<code>nnsreg_p</code>
<code>e(depvar)</code>	name of dependent variable
<code>e(indepvar)</code>	name of independent variable
<code>e(method)</code>	<code>connect</code> or <code>ols</code>
<code>e(model)</code>	connected or least-squares model label
<code>e(noise)</code>	<code>mean</code> or <code>median</code>
<code>e(partition)</code>	<code>xy</code> or <code>xonly</code>
<code>e(knots)</code>	interior knot locations, space separated
<code>e(vce)</code>	<code>ols</code> , <code>robust</code> , <code>cluster</code> , or <code>heuristic</code>
<code>e(vcetype)</code>	displayed covariance label, when applicable
<code>e(clustvar)</code>	cluster variable, when specified

Matrices

<code>e(b)</code>	coefficient vector
<code>e(beta)</code>	standardized coefficient vector
<code>e(V)</code>	variance-covariance matrix
<code>e(knotmat)</code>	interior knot locations, full precision
<code>e(segments)</code>	per-segment x range, slope, standard error, and n

Functions

<code>e(sample)</code>	marks estimation sample
------------------------	-------------------------

The coefficients in `e(b)` are named after the regressor (the first-segment slope) and `knot1`, `knot2`, and so on (the change in slope at each interior knot). `e(segments)` collects the per-segment slopes (cumulative sums of those coefficients) with delta-method standard errors, which is what the `slopes` option displays and what the `lincom` tests below compute one segment at a time. For `method(connect)`, `e(r2)` and `e(rmse)` are computed from raw squared residuals because the connected fit is not least squares, so its residuals need not average zero. The command also recomputes `e(rss)`, `e(mss)`, and `e(r2.a)` from the connected predictions and leaves the OLS-specific F and log-likelihood statistics missing. It labels the connected covariance as `e(vce)=heuristic` and the model as `e(model)=connect`, rather than retaining labels from the OLS scaffold. `margins` and `estat` are not supported, because the knot coefficients are spline-basis terms rather than variables in the dataset and the generated basis variables are temporary; `lincom` and `e(segments)` are the routes to marginal slopes. The `slopes` option

prints that table directly, one row per segment with its x range, slope, and standard error:

```
. use texas_education, clear
(Texas School District Operational Funding and Achievement (Simulated))
. nnsreg staar_pass spend_pupil, order(2) method(ols) noplot slopes
Nonlinear Nonparametric Statistics (NNS) Regression
Dependent variable:   staar_pass
Independent variable: spend_pupil
Method:               ols
Noise reduction:      mean
Partition:            xy
Order:                2
Number of obs:        250
R-squared:            0.9434
Root MSE:             4.2994
```

staar_pass	Coefficient	Std. err.	t	P> t	[95% conf. interval]	
spend_pupil	.3637342	.4431866	0.82	0.413	-.5092077	1.236676
knot1	7.75409	.6855804	11.31	0.000	6.403706	9.104473
knot2	-4.984399	.5169886	-9.64	0.000	-6.002708	-3.96609
knot3	-3.463223	.599266	-5.78	0.000	-4.643594	-2.282852
_cons	39.18212	3.16834	12.37	0.000	32.94146	45.42278

Per-segment slopes (cumulative), conditional on the NNS knots

segment	x range		slope	std. err.	n
1	5.01 to	8.28	0.3637	0.4432	58
2	8.28 to	11.8	8.1178	0.3223	60
3	11.8 to	16.4	3.1334	0.2532	67
4	16.4 to	20	-0.3298	0.4047	65

Standard errors are conditional on the data-derived knot locations.

4 Simulated comparisons

4.1 Design: three known curves, one seed, and 100 replications

All examples are simulated so that the true conditional mean is known, which allows each method to be scored on how well it recovers the curve that generated the data rather than only on how well it fits noise. The variable names borrow from Texas policy settings (school spending and a test-passing rate, water-system pipe age and water loss) to keep the examples concrete. The names are labels of convenience: all values are synthetic and stylized, and no claim in this article is about actual districts, assessments, or water systems.

The replication script `test_nnsreg.do`, included with this article, generates every number and figure shown here from fixed seeds. The script generates the following designs:

- **Education (smooth S-curve).** For 250 districts, spending per pupil x is uniform on $[5, 20]$ (thousands of dollars) and the true passing rate is $40 + 45/\{1 + \exp(-0.8[x - 11])\}$, an S-curve that is flat at low spending, rises through the middle, and plateaus near the top. Noise is normal with standard deviation 4.
- **Water (threshold).** For 300 systems, pipe age a is uniform on $[0, 80]$ years and the true loss rate is flat at 8% for $a \leq 35$ and $8 + 0.4(a - 35) + 0.01(a - 35)^2$ afterward, so the curve has a kink at age 35. Noise has standard deviation 2.5.
- **Clustered (peak and dip).** For 260 units, the score s is drawn from a three-component normal mixture (means 22, 36, and 58 with standard deviations 3, 2.2, and 4; mixture weights 0.45, 0.30, and 0.25), so the support is dense in three regions and thin between them. The true curve is $18 + 0.18s + 10g_{24,35} - 16g_{35,18} + 8g_{58,50}$, where $g_{c,w} = \exp\{-(s - c)^2/w\}$, which produces a local peak near 24, a sharp dip near 35, and a partial recovery near 58. Noise has standard deviation 1.2. This design mirrors the clustered-data setting that Vinod and Viole (2018b) identify as favorable to partition-based fits.

Each design has a canonical dataset with a fixed seed (2036, 2037, and 2038), used for the figures and the single-run table. Because any single draw can flatter one method by luck, the script also redraws each design 100 times with fresh seeds recorded in the script, refits every method, and stores the results; section 4.5 reports that study.

In-sample R^2 and RMSE against the observed outcome measure conformity to the data, including its noise. RMSE against the true curve, available only because the data are simulated, measures recovery of the signal. The distinction matters: a method can score better on observed RMSE by chasing noise while scoring worse on the truth, and the clustered example below shows exactly that pattern. Table 1 collects the single-run diagnostics for all three examples; the discussion of each example refers to its rows.

4.2 Example 1: smoothers lead on a smooth S-curve; the spline stays close

The education design asks whether each method can represent a plateau without the analyst choosing a polynomial order or hand-picking knots. Figure 1 shows four of the fits, one per panel, so each curve can be read against the data without overplotting; `npregress series` appears in table 1 rather than in the figure. The linear fit misses both the flat start and the plateau, and its true-curve RMSE of 5.460 is six to seven times that of the smoothers and the NNS OLS fit. The kernel smoother tracks the S-curve closely (true-curve RMSE 0.923) and `npregress series` slightly closer (0.811). The NNS OLS spline is closest in this particular draw (0.784), though the Monte Carlo study below shows the two `npregress` fits ahead on average; the connected NNS fit is visibly more jagged (3.283) because its coefficients interpolate regression points instead of minimizing squared error.

The log shows how `nnsreg` posts its coefficients. The first row reports the first-segment slope, each `knot` row reports a slope change, and `estimates table` accepts

Table 1: Single-run diagnostics from the canonical simulated datasets. “Observed RMSE” scores each fit against the noisy outcome; “true RMSE” scores it against the known generating curve. Both RMSE columns divide by N (a root mean squared deviation), which differs slightly from $e(\text{rmse})$, whose divisor is the residual degrees of freedom. All three fixed-seed draws flatter the NNS OLS spline relative to its replication average; section 4.5 reports the averages across 100 replications.

Example	Model	R^2	Observed RMSE	True RMSE
Education	Linear regression	0.864	6.611	5.460
Education	npregress kernel	0.950	4.015	0.923
Education	npregress series	0.950	4.007	0.811
Education	NNS connect, order(2)	0.924	4.939	3.210
Education	NNS connect, order(3)	0.927	4.828	3.283
Education	NNS connect, order(3) noise(median)	0.938	4.460	2.350
Education	NNS OLS, order(3)	0.949	4.027	0.784
Water	Linear regression	0.757	5.502	5.121
Water	npregress kernel	0.962	2.181	0.865
Water	npregress series	0.953	2.418	0.401
Water	NNS connect, order(2)	0.945	2.622	1.132
Water	NNS connect, order(3)	0.938	2.776	1.460
Water	NNS connect, order(3) noise(median)	0.929	2.979	1.804
Water	NNS OLS, order(3)	0.954	2.389	0.372
Water	NNS OLS, order(3) partition(xonly)	0.949	2.522	0.883
Clustered	Linear regression	0.042	9.234	9.219
Clustered	npregress kernel	0.985	1.162	0.484
Clustered	npregress series	0.985	1.138	0.392
Clustered	NNS connect, order(3)	0.976	1.468	0.999
Clustered	NNS OLS, order(3)	0.984	1.210	0.492
Clustered	NNS OLS, order(4)	0.987	1.078	0.452
Clustered	NNS OLS, order(5)	0.989	0.987	0.596

the stored results.

```
. use texas_education, clear
(Texas School District Operational Funding and Achievement (Simulated))
. nnsreg staar_pass spend_pupil, order(3) method(connect) noplot
Nonlinear Nonparametric Statistics (NNS) Regression
Dependent variable:  staar_pass
Independent variable: spend_pupil
Method:              connect
Noise reduction:     mean
Partition:           xy
Order:               3
Number of obs:       250
R-squared:            0.9272
Root MSE:            4.9380
```

staar_pass	Heuristic reference				
	Coefficient	std. err.	t	P> t	[95% conf. interval]




Figure 1: Education example, one fitted method per panel (orange) against the same simulated data (gray dots) and true curve (black). The linear model misses the plateau. The kernel smoother and the NNS OLS spline both recover the S-shape, with true-curve RMSE 0.923 and 0.784 in this draw; the connected NNS fit is more jagged because it interpolates regression points.

spend_pupil	6.414835	2.065966	3.11	0.002	2.345007	10.48466
knot1	-38.42267	12.13469	-3.17	0.002	-62.32728	-14.51806
knot2	33.71038	11.14115	3.03	0.003	11.763	55.65776
knot3	5.758648	1.765085	3.26	0.001	2.281537	9.235759
knot4	1.668066	7.469167	0.22	0.823	-13.04574	16.38187
knot5	6.868115	13.09004	0.52	0.600	-18.91846	32.65469
knot6	-10.78198	7.699945	-1.40	0.163	-25.9504	4.386446
knot7	-8.614286	1.875776	-4.59	0.000	-12.30945	-4.919122
knot8	9.589458	1.80015	5.33	0.000	6.043272	13.13564
knot9	-5.920695	1.735984	-3.41	0.001	-9.340479	-2.500911
_cons	6.99154	11.78148	0.59	0.553	-16.21726	30.20034

note: connect-fit standard errors are heuristic references from a rescaled OLS
> spline covariance; they are not estimator-specific tests.

```
. estimates store edu_connect
. quietly nnsreg staar_pass spend_pupil, order(3) method(ols) noplot
. estimates store edu_ols
```

```
. estimates table edu_connect edu_ols, b(%9.3f) se(%9.3f) stats(r2 rmse N)
```

Variable	edu_con~t	edu_ols
spend_pupil	6.415	1.597
	2.066	1.723
knot1	-38.423	-9.011
	12.135	10.122
knot2	33.710	10.028
	11.141	9.293
knot3	5.759	4.362
	1.765	1.472
knot4	1.668	5.818
	7.469	6.230
knot5	6.868	-3.221
	13.090	10.919
knot6	-10.782	-4.405
	7.700	6.423
knot7	-8.614	-4.087
	1.876	1.565
knot8	9.589	-0.175
	1.800	1.502
knot9	-5.921	-0.842
	1.736	1.448
_cons	6.992	32.086
	11.781	9.828
r2	0.927	0.949
rmse	4.938	4.119
N	250	250

Legend: b/se

Figure 2 plots the stored coefficients with `coefplot` (Jann 2014), with the knot locations from `e(knots)` written into the coefficient labels. The two knots near spending of 6.5 and 6.7 bound a short range with no observations, so the fitted line there is interpolation across a gap rather than evidence of a local change. The associated wide intervals carry little information. The readable pattern is in the interior: positive slope changes where the curve is rising are offset by negative changes higher up the spending range, and that offset is what flattens the fit into the plateau.

Because the plateau is the substantive claim, the fitted terminal slope provides a direct numerical summary. It equals the first-segment coefficient plus every knot coefficient, which is one `lincom` statement:

```
. estimates restore edu_ols
(results edu_ols are active now)
. lincom spend_pupil + knot1 + knot2 + knot3 + knot4 + knot5 + knot6 ///
> + knot7 + knot8 + knot9
( 1) spend_pupil + knot1 + knot2 + knot3 + knot4 + knot5 + knot6 + knot7 +
    knot8 + knot9 = 0
```

staar_pass	Coefficient	Std. err.	t	P> t	[95% conf. interval]
(1)	.0645397	.7393695	0.09	0.931	-1.391973 1.521053




Figure 2: Stored slope-change components for the education example, labeled with knot locations from `e(knots)`. The 6.5–6.7 range contains no observations, so its fitted movement is interpolation across a gap. Over the supported interior, positive slope changes where the curve is rising are offset by negative changes higher up the spending range, producing the plateau.

The estimated terminal slope is 0.06 percentage points per thousand dollars (95% conditional interval -1.39 to 1.52), consistent with the plateau built into the generating curve. The true secant slope over the fitted terminal range is 0.08. These intervals are conditional on the estimated knot locations, so I read them as descriptive support rather than exact inference. Still, the question “is the top of the curve flat?” becomes a one-line test on stored coefficients rather than a visual judgment.

4.3 Example 2: on a threshold curve, the estimated knots land near the true change point

The water design has a kink: the true loss rate is flat until pipe age 35 and rises afterward. Figure 3 shows the same four-panel comparison, and table 1 again covers `npregress series`. The kernel smoother, the series fit, and the NNS OLS spline all recover the flat-then-rising shape (true-curve RMSE 0.865, 0.401, and 0.372 in this draw); the linear fit averages the two regimes and represents neither.




Figure 3: Water example, one fitted method per panel. The true curve is flat until pipe age 35 and rises afterward. The kernel smoother and the NNS OLS spline both recover the shape (true-curve RMSE 0.865 and 0.372); the linear fit averages the flat and rising regimes and represents neither.

The segmented fit also reports where its slope changes. Figure 4 overlays the interior knots stored in `e(knots)` over the data: several knots land near the true change point at age 35, and the fitted slope changes concentrate there. Two `lincom` statements quantify the same pattern, one for the first segment and one for the last:

```
. use texas_water, clear
(Texas Municipal Water Infrastructure and Water Loss (Simulated))
. quietly nnsreg water_loss pipe_age, order(3) method(ols) noplot
. lincom pipe_age
( 1) pipe_age = 0
```

	water_loss	Coefficient	Std. err.	t	P> t	[95% conf. interval]
(1)		-.1292191	.1497085	-0.86	0.389	-.4238763 .1654382

```
. lincom pipe_age + knot1 + knot2 + knot3 + knot4 + knot5 + knot6 ///
```

```
> + knot7 + knot8 + knot9
( 1) pipe_age + knot1 + knot2 + knot3 + knot4 + knot5 + knot6 + knot7 + knot8
+ knot9 = 0
```

water_loss	Coefficient	Std. err.	t	P> t	[95% conf. interval]	
(1)	1.209637	.0946187	12.78	0.000	1.023408	1.395866

The fitted slope among the youngest systems is -0.13 percentage points per year (95% conditional interval -0.42 to 0.17), and the interval includes the true value of zero. The terminal slope is 1.21 (95% conditional interval 1.02 to 1.40), close to the true secant slope of 1.19 over the fitted terminal range. Table 1 also reports the `partition(xonly)` variant for this example (true RMSE 0.883): the vertical-band partition is less accurate here than the quadrant default but reproduces the same substantive contrast. I would report this descriptive pattern only when the two partition styles agree in this way.




Figure 4: The NNS OLS fit for the water example with its interior knots (dashed gray lines, from `e(knots)`) and the true change point at age 35 (navy line). Several knots land near the true kink, and the fitted curve is flat to the left of it and rising to the right, matching the slope-contrast tests in the text.

4.4 Example 3: clustered support favors the NNS least-squares fit

The clustered design is the setting the NNS literature identifies as favorable (Vinod and Viole 2018b): observations are dense in three regions with thin support between them, and the true curve has a sharp dip. A kernel smoother must choose one bandwidth

profile for data whose local density varies sharply across the support; the NNS partitions instead adapt to where the observations sit.




Figure 5: Clustered example, one fitted method per panel. Observations concentrate in three regions, with a sharp dip in the true curve near the middle cluster. In this draw the kernel smoother and the NNS OLS spline are nearly tied (true-curve RMSE 0.484 and 0.492); the Monte Carlo study shows that across repeated draws the segmented fit is more accurate than the kernel in this design (table 2).

Figure 5 shows the single-run comparison, and it also shows why single runs mislead: in this particular draw the kernel smoother (0.484) and the NNS OLS spline (0.492) are nearly tied, and **npregress series**, again in table 1 rather than the figure, edges both (0.392). The Monte Carlo study in section 4.5 shows that the series estimator's average advantage comes with much larger draw-to-draw variability. Table 1 adds the **order()** progression, which previews the overfitting tradeoff: moving from **order(3)** to **order(4)** improves both observed and true RMSE, but **order(5)** improves observed RMSE (0.987) while worsening true RMSE (0.596). Extra segments eventually fit noise, and only the true-curve metric reveals it.

4.5 The Monte Carlo study: smoothers lead on smooth curves; NNS OLS is steadier on clustered support

A comparison based on one simulated dataset is weak evidence, so the replication script redraws every design 100 times (fresh regressor values and fresh noise each time), refits linear regression, `npregress kernel`, `npregress series`, and three NNS specifications, and scores each fit against the true curve. Table 2 reports the mean and standard deviation of true-curve RMSE across replications and the share of replications in which each competitor beats the kernel smoother. Figure 6 shows the full distributions. Because every method is refit on the same draw, the per-design comparisons are paired, which reduces noise in the differences between methods. Every estimator fit all 300 design-replications without error.

Table 2: Monte Carlo summary across 100 replications per design. Cells show the mean of RMSE against the true curve, its standard deviation across replications, and the share of replications in which the model beats `npregress kernel` on that metric. The `npregress` estimators lead both smooth designs; in the clustered design the NNS OLS fits trade a little average accuracy for much lower between-replication variability.

Design	Model	Mean true RMSE	SD	Beats kernel
Education	Linear regression	5.512	0.169	0.00
Education	<code>npregress kernel</code>	0.802	0.194	–
Education	<code>npregress series</code>	0.749	0.219	0.62
Education	NNS connect, <code>order(3)</code>	3.168	0.596	0.00
Education	NNS OLS, <code>order(3)</code>	0.887	0.180	0.28
Education	NNS OLS, <code>order(4)</code>	1.290	0.180	0.00
Water	Linear regression	5.116	0.132	0.00
Water	<code>npregress kernel</code>	0.426	0.104	–
Water	<code>npregress series</code>	0.416	0.082	0.52
Water	NNS connect, <code>order(3)</code>	1.531	0.273	0.00
Water	NNS OLS, <code>order(3)</code>	0.483	0.095	0.23
Water	NNS OLS, <code>order(4)</code>	0.784	0.100	0.00
Clustered	Linear regression	9.172	0.301	0.00
Clustered	<code>npregress kernel</code>	0.716	0.233	–
Clustered	<code>npregress series</code>	0.467	0.333	0.89
Clustered	NNS connect, <code>order(3)</code>	1.114	0.252	0.10
Clustered	NNS OLS, <code>order(3)</code>	0.540	0.091	0.67
Clustered	NNS OLS, <code>order(4)</code>	0.496	0.067	0.80

The results split by the shape of the curve, and the two `npregress` estimators behave alike. In the two smooth designs both smoothers are more accurate on average than any NNS specification: in the education design the kernel averages 0.802 and series 0.749, against 0.887 for the best NNS fit, and in the water design the three are close (0.426, 0.416, and 0.483). The NNS OLS spline beats the kernel in about a quarter of the smooth-design replications, and series beats it in 52% (water) and 62% (education) of replications, so the two `npregress` flavors are near-substitutes here. Vinod and




Figure 6: Distribution of true-curve RMSE across 100 Monte Carlo replications, by design and method (linear regression omitted; its values sit far above every panel and appear in table 2). The two **npregress** estimators are most accurate in the smooth designs. In the clustered design **npregress series** has the lowest average error but the widest spread, while the NNS OLS fits are nearly as accurate on average and have much narrower distributions. The vertical scale differs by design, and RMSE is in the outcome’s units.

Viole (2018b) report that in-sample fit is similar for NNS and kernel regression, with differences they call minor, and these smooth designs are consistent with that reading.

The clustered design separates average accuracy from variability across replications. **npregress series** has the lowest mean RMSE (0.467), while NNS OLS at orders 4 and 3 averages 0.496 and 0.540; all three are below the kernel’s 0.716. Series has the widest RMSE distribution in this design, with a standard deviation of 0.333 across replications, larger than the kernel’s 0.233 and roughly four to five times the NNS OLS spread of 0.07 to 0.09. The summary does not identify the mechanism, although thin support gives a flexible series fit room to move between clusters. The NNS OLS fits reduce the between-replication standard deviation to roughly one-fifth to one-quarter of the series value. This is the setting the NNS literature singles out as favorable to partition-based fits (Vinod and Viole 2018b), argued there from an example rather than a benchmark; in this design, the NNS OLS advantage is lower between-replication variability on clustered support.

The Monte Carlo results show why the fixed-seed examples should not determine the ranking. The water draw in table 1 shows the NNS spline beating the kernel (0.372 versus 0.865), and the education draw does the same (0.784 versus 0.923), yet across 100 draws the kernel is slightly better on average in both smooth designs. In the clustered draw, the NNS OLS spline and kernel are nearly tied; across replications, both NNS OLS fits beat the kernel on average, while the series fit has the lowest mean error. All three fixed-seed draws happen to flatter the spline relative to its replication average. I

report both layers because single simulated examples, including mine, are illustrations rather than evidence.

Finally, the `order(4)` rows are a caution against treating more segments as better. The higher order helps only in the clustered design, where the true curve has sharp local features; in both smooth designs it is worse than `order(3)` because the extra segments track noise.

Two further checks probe how far these results carry. Because the article reports conditional intervals as a deliverable, I measure their calibration directly. Across the same 100 replications, table 3 reports how often the nominal 95% conditional interval for the terminal-segment slope contains the true secant slope of the generating curve over that segment. The least-squares intervals cover at close to the nominal rate, 0.94 to 0.98 across the three designs, so reading them as descriptive summaries is reasonable even though they ignore the selection of the knots; the connected-fit intervals are less consistent, 0.80 to 0.99, which is why I call them heuristic references rather than tests. The clustered variance advantage is also not tied to one noise level. Repeating the clustered design with the noise standard deviation raised from 1.2 to 2.0, the NNS OLS spline again matches `npregress series` on mean true-curve RMSE, 0.75 for both, with about a fifth of its replication spread (standard deviation 0.10 against 0.51); a linear spline with six rank-placed percentile knots is far less accurate (mean 1.48). Placing knots from the data rather than at fixed ranks helps even when the truth is known. Both checks run in `robustness.do`.

Table 3: Empirical coverage of the nominal 95% conditional interval for the terminal-segment slope, across 100 replications per design. The target is the true secant slope of the generating curve over the fitted terminal segment.

Design	<code>method(ols)</code>	<code>method(connect)</code>
Education	0.94	0.80
Water	0.94	0.99
Clustered	0.98	0.99

4.6 The reported shapes survive order and center changes; extra depth tracks noise

Because `order()` and `noise()` are judgment calls, the practical question is whether the substantive reading survives reasonable variations. Figures 7 and 8 overlay connected NNS fits across `order(2)`, `order(3)`, and `order(3) noise(median)` for the two policy-flavored examples. In both, the shape that matters (the plateau in one, the flat-then-rising threshold in the other) appears in every variant, which is the pattern that would justify reporting it. Figure 9 runs the same exercise over `order()` in the clustered design and shows the failure mode instead: `order(5)` adds local movement that improves observed fit while drifting from the true curve, the same pattern as in table 1. When a finding disappears between `order(2)` and `order(3)`, or between mean and median

centers, report it as tentative.




Figure 7: Education example: connected NNS fits across order and center choices. The S-shape and upper plateau appear in every variant, though the `order(3)` mean fit adds local movement near spending of 6.5 and 16. The shared shape is the stability pattern that supports reporting the plateau.

5 A real-data illustration

The simulated designs let me score how well each method recovers a curve I built, but they are stylized by construction. To watch `nnsreg` on data I did not design, I turn to the motorcycle-impact benchmark: accelerometer readings from a crash-helmet experiment, measured against time after impact. The data come from Silverman (1985), are reprinted in Fan and Gijbels (1996), and ship with Stata Press as `motorcycle.dta` (the dataset in the [R] `lpoly` examples). They are an intended stress case for `nnsreg`: the regressor is heavily tied and clustered (94 distinct times among 133 readings), the mean response is flat before impact and then plunges into a trough and rebounds, and the noise grows with time. Because the true mean is unknown here, I judge accuracy by held-out error from repeated cross-validation, the real-data stand-in for the true-curve RMSE of section 4.

5.1 On the impact trough, the segmented fit tracks the smoothers

Figure 10 shows four fits, one per panel. The linear fit cannot represent a curve that reverses twice (in-sample RMSE 45.98 against the observed acceleration). The kernel




Figure 8: Water example: connected NNS fits across order and center choices. Every variant shows the flat region below age 35 and the rise after it; **order(2)** is smoother, and **order(3)** allows more local movement through the flat region and near the threshold.

smoother and the NNS OLS spline both trace the flat pre-impact run, the trough, and the rebound, with in-sample RMSE 21.74 and 21.47; **npregress series** is close behind at 21.72. The connected NNS fit is visibly rougher (30.64), the same jaggedness the simulated examples showed, because its coefficients interpolate the regression points rather than minimizing squared error.

5.2 The knots concentrate in the impact trough

The interpretive payoff is the same as in the simulated examples: the fit decomposes the relationship into ranges that an analyst can inspect. Figure 11 shows the fitted curve and 13 interior knots above the slope assigned to each resulting range. Many knots fall in the trough between about 14 and 32 milliseconds. Their locations reflect both the joint regressor-outcome partition and the empirical support, so they flag that range for inspection rather than estimate a threshold. The per-segment slopes from the **slopes** option put numbers on the shape: the fitted slope is flat before impact (0.09 g per millisecond in the first segment, with a robust standard error of 0.30), plunges through the trough (−22.8 and −31.0 g per millisecond in the two segments spanning 14.5 to 17.7 milliseconds), and rebounds afterward (+25.6 near 24 to 26 milliseconds). Several ranges contain fewer than five observations, and the short interval from 45.7 to 45.9 milliseconds contains none; its slope is an interpolation between adjacent knots. The lower panel exposes those weakly supported ranges rather than hiding them. The displayed robust intervals describe uncertainty conditional on the selected knots.




Figure 9: Clustered example: NNS OLS fits at orders 3 through 5 against the kernel benchmark. Orders 3 and 4 recover the peak, dip, and recovery; **order(5)** adds movement that improves observed fit but drifts from the true curve, the overfitting pattern quantified in table 1.

5.3 Out of sample, the segmented fit and kernel have similar mean error

A fit that follows the data in sample may only be chasing noise, so table 4 reports held-out root mean squared error from repeated five-fold cross-validation (20 repeats, folds redrawn each time). The smallest and largest time are always kept in training, so each repeat scores 131 of 133 observations without extrapolation. Because folds are assigned by observation, tied time values can appear in both training and validation; the exercise estimates error for another reading at a partly familiar time, not at an unseen time value. The NNS OLS spline and the kernel smoother have nearly identical mean held-out RMSE, 24.55 and 24.52 (standard deviations 0.77 and 0.79 across the 20 fold assignments). The series fit has the lowest mean, 24.08, with a standard deviation of 1.29. The median-centered NNS OLS fit is slightly behind at 24.61. The connected fit is worse and less steady (28.02, standard deviation 1.76), as expected of the descriptive rather than the accuracy-competitive variant. The two smoothers and all three segmented specifications have lower mean error than a restricted cubic spline with five percentile knots (35.62), whose rank-placed knots miss the trough; every flexible fit also has lower mean error than the linear fit (46.53).




Figure 10: Motorcycle data, one fitted method per panel (dark line) against the 133 readings (gray dots). The linear fit cannot represent a curve that reverses; the kernel smoother and the NNS OLS spline both trace the flat run, the impact trough, and the rebound (in-sample RMSE 21.74 and 21.47).




Figure 11: Motorcycle segment diagnostic. The upper panel shows the NNS OLS fit and its 13 interior knots. The lower panel shows the slope over each range, a robust 95% interval conditional on the selected knots, and the number of observations in the range. Dashed gray ranges contain fewer than five observations; one contains none and represents interpolation between knots. Large slopes and dense knot placement flag ranges for follow-up, not estimated thresholds.

Table 4: Motorcycle data: fit quality in sample and out of sample. In-sample RMSE is against the observed acceleration; held-out RMSE is the mean (and standard deviation) across 20 repeats of five-fold cross-validation, lower is better. The endpoint rule means that each repeat scores 131 of 133 readings. The restricted-cubic-spline basis uses five percentile knots fixed from the full regressor distribution before the folds. **npregress series** is scored with **predict, tolerance(100)**; **nnsreg** rebuilds the basis from **e(knots)** and scores held-out rows directly.

Method	In-sample RMSE	Held-out RMSE	SD
Linear regression	45.98	46.53	0.32
Restricted cubic spline, 5 knots	—	35.62	0.66
npregress kernel	21.74	24.52	0.79
npregress series	21.72	24.08	1.29
NNS connect, order(3)	30.64	28.02	1.76
NNS OLS, order(3)	21.47	24.55	0.77
NNS OLS, order(3) noise(median)	21.74	24.61	0.93

The motorcycle benchmark adds one held-out comparison but does not reproduce the simulation's comparison of between-replication variability. The NNS OLS series fit has the lowest mean held-out error; the NNS OLS spline and kernel are close behind and nearly identical to each other. The NNS fit also reports range-specific slopes and support counts and scores held-out rows directly from its stored spline without a tolerance option, a practical convenience for a rerunnable workflow.

6 Practical guidance

The comparisons suggest a division of labor. Consider `npregress kernel` or `npregress series` when the goal is the most accurate estimate of a smooth conditional mean and the sample covers the range of interest evenly; the two were the most accurate methods in both smooth designs, and `npregress series` is the closest competitor to `nnsreg` because it too posts a data-driven spline as estimation results. When the goal is instead formal inference on a small number of break locations, a segmented or threshold estimator is the right tool, not `nnsreg`.

Consider `nnsreg` when the deliverable is a segment-level summary: named slope coefficients, knot locations in `e(knots)`, and `lincom`-computable contrasts such as the plateau and threshold summaries above, none of which `npregress` reports. The command is also worth comparing when the regressor's support is clustered or gapped. In the clustered simulation studied here, `npregress series` can reach a low average error but swings widely from draw to draw, while `method(ols)` delivers about the same average accuracy with a fraction of the draw-to-draw spread. The default connected fit is the descriptive, partition-faithful variant; `method(ols)` is the accuracy-competitive version in these comparisons. Comparing the two fits tests whether a reported feature depends on one tuning approach. A bandwidth-free `nnsreg` fit offers an alternative representation alongside a kernel or series fit. Agreement strengthens the descriptive reading: when a kernel smoother and the NNS partition both show a plateau or a trough, the feature is less likely an artifact of one smoothing choice. Disagreement makes the shape tentative. The command replaces bandwidth selection with choices about `order()`, `minobs()`, and partitioning, which the sensitivity grid of section 4.6 evaluates. Table 5 lays the contrasts side by side.

Results from `nnsreg` need the same scrutiny as any data-driven fit. Fit a small grid of `order()` values and both centers, and report whether the shape survives, as in section 4.6. And when prediction accuracy is the question, compare methods on a common metric that penalizes noise-chasing, such as the held-out cross-validation error of the motorcycle example in section 5; the true-curve RMSE used with the simulated designs is the same idea when the generating curve is known.

An unexpected local slope is a prompt to inspect the data, not a finding by itself. Start with the range's observation count in `e(segments)` and look for gaps, outliers, coding problems, or changes in subgroup composition. Then compare `partition(xy)` with `partition(xonly)`, vary `order()`, `minobs()`, and the center, and place a kernel or series fit beside the segmented fit. A well-supported feature that survives these checks

Table 5: What each command reports. **nnsreg** scores held-out observations from stored results and derives its knots from the NNS partition. **npregress series** scores new rows when **tolerance()** or **atsample** is specified. **mkspline** uses locations supplied by the user or placed by rank.

	nnsreg	npregress kernel	npregress series	mkspline + regress lpoly	
Posts e-class results	yes	yes	yes	yes	no
No bandwidth to choose	yes	no	yes	yes	no
Named segment-slope contrasts	yes	no	no	yes	no
predict out of sample	yes	no	with option	yes	no
Data-derived knot locations	NNS partition	no	automatic	user/rank	no

is a descriptive pattern worth investigating further; it is not a causal threshold.

Computation is modest in these examples: each individual fit completes in well under a second, and the full 100-replication Monte Carlo study, including 300 **npregress kernel** fits, completes in under three minutes on the machine used for this article.

7 Conclusion

nnsreg adapts the NNS regression method of Viole and Nawrocki (2013) to Stata and posts its output through Stata’s standard estimation interface. The simulations delimit the command’s role: the **npregress** estimators recover smooth curves more accurately on average, while on the clustered, sharply featured benchmark studied here the NNS OLS spline nearly matches the most accurate smoother on average and has much lower between-replication variability. In every design it yields a range-specific summary with reportable slopes, support counts, and explicit uncertainty caveats that a smooth curve does not.

Future extensions could add multivariate regressors and weights, resampling-based uncertainty summaries, and automated order selection. The current interval estimates are conditional on data-derived knots. Because the number and location of knots change from resample to resample, coefficient-wise bootstrap estimates do not align naturally; a useful upgrade would repeat the partitioning and summarize quantities every resample shares, such as the fitted curve at fixed grid points or slopes at fixed quantiles of the regressor. The command also asks users to choose **order()** by comparison rather than automating the choice; the dependence-based order selection in the R package (Viole 2026) and cross-validation over a small **order()** grid offer two possible designs for a future **order(auto)** option. Each extension requires validation before it can support additional automation.

8 Acknowledgments

The NNS framework is the work of Fred Viole and David Nawrocki, and the cluster-based regression variant adapted here was developed by Hrishikesh Vinod and Fred Viole. This article adapts their method to Stata; any errors in the adaptation are mine.

9 References

- Bawa, V. S. 1975. Optimal Rules for Ordering Uncertain Prospects. *Journal of Financial Economics* 2(1): 95–121.
- Cattaneo, M. D., R. K. Crump, M. H. Farrell, and Y. Feng. 2024. On Binscatter. *American Economic Review* 114(5): 1488–1514.
- . 2025. Binscatter Regressions. *Stata Journal* 25(1): 3–50.
- Fan, J., and I. Gijbels. 1996. *Local Polynomial Modelling and Its Applications*. London: Chapman & Hall.
- Fishburn, P. C. 1977. Mean-Risk Analysis with Risk Associated with Below-Target Returns. *American Economic Review* 67(2): 116–126. <https://www.jstor.org/stable/1807225>.
- Friedman, J. H. 1991. Multivariate Adaptive Regression Splines (with discussion). *Annals of Statistics* 19(1): 1–67.
- Hansen, B. E. 2017. Regression Kink with an Unknown Threshold. *Journal of Business & Economic Statistics* 35(2): 228–240.
- Hayfield, T., and J. S. Racine. 2008. Nonparametric Econometrics: The np Package. *Journal of Statistical Software* 27(5): 1–32.
- Jann, B. 2014. Plotting Regression Coefficients and Other Estimates. *Stata Journal* 14(4): 708–737.
- Kim, H.-J., M. P. Fay, E. J. Feuer, and D. N. Midthune. 2000. Permutation Tests for Joinpoint Regression with Applications to Cancer Rates. *Statistics in Medicine* 19(3): 335–351.
- Muggeo, V. M. R. 2003. Estimating Regression Models with Unknown Break-Points. *Statistics in Medicine* 22(19): 3055–3071.
- Nawrocki, D., and F. Viole. 2024. Estimation Error and Partial Moments. *International Review of Financial Analysis* 95: 103443.
- Newson, R. 2006. Confidence Intervals for Rank Statistics: Somers’ D and Extensions. *Stata Journal* 6(3): 309–334.
- Orsini, N., and S. Greenland. 2011. A Procedure to Tabulate and Plot Results after Flexible Modeling of a Quantitative Covariate. *Stata Journal* 11(1): 1–29.

- Royston, P. 2013. marginscontplot: Plotting the Marginal Effects of Continuous Predictors. *Stata Journal* 13(3): 510–527.
- Royston, P., and D. G. Altman. 1994. Regression Using Fractional Polynomials of Continuous Covariates: Parsimonious Parametric Modelling. *Journal of the Royal Statistical Society, Series C (Applied Statistics)* 43(3): 429–467.
- Ruppert, D., M. P. Wand, and R. J. Carroll. 2003. *Semiparametric Regression*. Cambridge: Cambridge University Press.
- Silverman, B. W. 1985. Some Aspects of the Spline Smoothing Approach to Non-Parametric Regression Curve Fitting (with discussion). *Journal of the Royal Statistical Society, Series B (Methodological)* 47(1): 1–52.
- StataCorp. 2025. *Stata Statistical Software: Release 19*. StataCorp LLC, College Station, TX.
- Vinod, H. D., and F. Viole. 2018a. Clustering and Curve Fitting by Line Segments. Preprints.org working paper 2018010090.
- . 2018b. Nonparametric Regression Using Clusters. *Computational Economics* 52(4): 1317–1334.
- Viole, F. 2026. *NNS: Nonlinear Nonparametric Statistics*. Comprehensive R Archive Network. R package version 13.1, available at <https://CRAN.R-project.org/package=NNS>.
- Viole, F., and D. Nawrocki. 2012a. Cumulative Distribution Functions and UPM/LPM Analysis. SSRN working paper 2148482. <https://ssrn.com/abstract=2148482>.
- . 2012b. Deriving Nonlinear Correlation Coefficients from Partial Moments. SSRN working paper 2148522. <https://ssrn.com/abstract=2148522>.
- . 2013. *Nonlinear Nonparametric Statistics: Using Partial Moments*. CreateSpace Independent Publishing Platform.

About the author

Eric A. Booth is a Senior Researcher at Texas 2036, a nonpartisan public-policy think tank in Austin, TX. His applied work focuses on quantitative policy analysis methods, education policy and program evaluation, and reproducible Stata workflows for applied research.